

Claude Code for Biomolecular Simulation & Computational Chemistry

From First Principles to the BioKT Lab Configuration

David De Sancho

June 6, 2026

University of the Basque Country (UPV/EHU) & DIPIC

Prelude

Live Demo



Prelude: /dft-homo-lumo on Guanidinium

What just happened

- One prompt: molecule name and charge
- Claude read a skill file encoding the full protocol
- Generated 3D geometry → DFT optimisation → SCF
- Extracted HOMO/LUMO energies and cube files
- Notebook visualised the frontier orbitals

Key numbers

	Guan. (+1)	Guan. (o)
HOMO (eV)	−15.08	−8.15
LUMO (eV)	−3.51	+1.47
Gap (eV)	11.57	9.63

Method: CAM-B3LYP / def2-TZVP



A Qualitative Shift in Research Practice

The jagged frontier

AI is superhuman at some tasks and surprisingly poor at others. Knowing which is which requires genuine domain expertise — not prompting skill.

Expertise is rewarded, not replaced

The productivity gains from AI go disproportionately to those who already know their field deeply — because they know when to trust the output and when to push back.

The bottleneck shifts

AI agents can work on many tasks in parallel. The limiting factor is no longer execution — it is the quality of human judgment directing the work.

The tool does not determine the outcome.
The expertise behind it does.



Outline

Part I — General Concepts

- What Claude Code is (and is not)
- Installation, interfaces, operating modes
- CLAUDE.md, skills, memory, sessions
- Cost, privacy, and honest caveats

Part II — The BioKT Lab

- Scientific standards and compute infrastructure
- Skills library: QM, MD, enhanced sampling
- Daily workflow in practice

Closing

- What you can do starting Monday
- A caution on outsourcing judgment

Questions welcome throughout — or at the end.



Part I

General Concepts



What Is Claude Code?

What it IS

- An **agentic AI coding assistant**
- Lives in your terminal or IDE
- Reads and writes files, runs shell commands
- Calls external tools and browses the web
- Operates in a loop: plan → act → observe & validate
- Controlled by you at each step (or autonomously)

What it is NOT

- A chatbot (it *does* things, not just talks)
- An autocomplete plugin
- A one-shot code generator
- A replacement for understanding your domain

Key insight

You describe *intent*; Claude figures out the steps, executes them, and reports back.



Installation & Setup

Terminal (Mac / Linux)

- **Install:** `curl -fsSL https://claude.ai/install.sh | bash`
- **Launch & auth:** `claude` — browser opens for login
- Subscription or API key required
- Auto-updates in the background

macOS alternative: `brew install --cask claude-code`

VS Code Extension

- Extensions panel (Ctrl+Shift+X)
- Search **Claude Code** — publisher: Anthropic
- Provides: file tree, real-time diffs, inline edit review

First steps after install

- Edit `~/.claude/CLAUDE.md` — your standing brief
- Run `/init` inside a project to auto-generate a local `CLAUDE.md`



Cost & Privacy

Cost

- **Claude Pro** — ~\$20/month; sufficient for most individual use
- **Claude Max** — higher usage limits; better for heavy or autonomous workflows
- **API / Console** — pay per token; useful for scripting and batch tasks
- No per-seat cost for the `claude` CLI itself

Privacy

- Anthropic does **not** use your conversations to train models by default
- Treat unpublished data with caution — avoid pasting novel sequences, unreleased structures, or confidential results into the session
- Work with file paths and structure; minimise raw scientific data in the conversation window

Rule of thumb

Treat Claude as an unaffiliated third-party service: fine for published data, cautious for unpublished results, avoid for anything under IP protection or confidentiality obligations.



Two Ways to Run Claude Code

Terminal

\$ claude

- Text-only interface
- Fast and lightweight
- Steeper initial barrier
- Ideal for scripting and quick tasks

IDE Extension (VS Code)

Claude Code extension

- File tree and real-time diffs
- Review edits before accepting
- Lower barrier for new users
- Ideal for extended coding sessions
- Comfortable for remote work via SSH on different machines



Four Operating Modes

Plan

- Read-only exploration
- No files are changed
- Think before acting

Default

- Asks before every action
- Safest starting point
- Full control

Accept Edits

- File edits auto-applied
- Shell commands still confirmed
- Good middle ground

Auto

- Fully autonomous
- Runs commands freely
- Fastest; use with care

Shift+Tab cycles between modes

| Start cautious, gain speed through trust



CLAUDE.md — Your Instruction File

What it is

- A Markdown file loaded into *every* session
- The “system prompt you write”
- Global: `~/ .claude/CLAUDE.md`
- Local: `<project>/CLAUDE.md` (project-specific)
- Changes Claude’s behaviour without touching the model

Each session begins with a fresh context window — CLAUDE.md is how you carry knowledge across sessions.

What to put in it

- **Role & identity** — who you are, your domain
- **Standards** — correctness rules, safety constraints
- **Work preferences** — how to communicate, what to ask
- **Infrastructure** — machines, paths, software stack
- **Skill catalogue** — which /skills exist and when to use them
- **Project structure** — layout, git conventions



Skills & Slash Commands

What they are

- Reusable workflows stored as SKILL.md files
- Triggered by typing /skill-name in the prompt
- Claude reads the instructions and executes them
- Loaded at startup — restart after adding new skills

Scope

- **Global:** ~/.claude/skills/ — everywhere
- **Local:** .claude/skills/ — project only; in git

Not all /commands are skills — /clear, /compact, /help are built-in CLI commands that execute directly, not through Claude.



Example

```
# In the Claude prompt:  
/review-plan  
  
# Claude stress-tests your plan:  
# - Identifies blind spots  
# - Checks best practices  
# - Flags wishful thinking  
# - Proposes alternatives
```

Why they matter

- Encode expertise *once*, reuse forever
- Share across machines and team members
- Build a personalised AI toolkit

Memory System

What is memory?

Each Claude session starts from scratch — it remembers nothing from previous conversations. The **memory system** solves this by writing persistent Markdown files that are loaded into every new session, carrying knowledge across conversations and machines.

Four memory types

Type	Stores
user	Role, expertise, preferences
feedback	Corrections & confirmed approaches
project	Goals, deadlines, decisions
reference	Where to find external resources

How it works

Files in

`~/.claude/projects/*/memory/`

`MEMORY.md` index loaded in every session

Claude saves memories proactively and on request

Prevents repeating the same corrections

Cross-session and cross-machine continuity



Session Management

What is context?

Context is everything Claude can currently “see”: the conversation history, file contents read, tool outputs, and loaded instructions (CLAUDE.md, memory). The model can only act on what is in context — nothing else exists for it.

Why it matters

- Sessions grow; model attention dilutes
- Longer session $\neq$ better outcome
- Each turn becomes more expensive over time
- Strategic breaks produce cleaner results

Key principle

Compact at task boundaries — not in a crisis.

Core commands

- `/compact <hint>` — summarise & continue
Hint tells Claude what to keep; be specific
- `/clear` — empty conversation, fresh start
Best for switching to an unrelated task
- `/rewind (Esc Esc)` — restore a prior state
Keep file edits but drop the discussion, or vice versa



Settings & Hooks

settings.json key options

model — which Claude model to use
(sonnet, opus, haiku)

permissions.allow — commands that run
without a confirmation prompt

hooks — shell prompts triggered on events
(e.g. PreCompact, PostToolCall)

statusLine — dynamic terminal status display

MCP servers — connect Claude to external tools
databases, APIs, local services; worth exploring once
comfortable with the basics

Global: ~/.claude/settings.json

Local: settings.local.json

Hook example

```
"hooks": {  
  "PreCompact": [{  
    "type": "prompt",  
    "prompt": "Run /ctxt save  
to checkpoint the session  
before compaction."  
  }]  
}
```

Before context is compressed, Claude
automatically saves the session state.



Honest Caveats

What Claude does well

- Routine, well-documented protocols
- Boilerplate code, file I/O, scripting
- Aggregating and restructuring information
- Catching obvious errors in structured output

Where it falls short

- Domain-specific edge cases with sparse training data
- Subtle physical reasoning (e.g. rare force field combinations, unusual collective variables)
- Knowledge of *your* system, your history, your failures
- It can be confidently wrong

The bottom line

Claude produces drafts. Validation is always your responsibility.




Part II

The BioKT Lab Configuration



BioKT Setup at a Glance



“Correctness over speed”

Simulation work goes on the scientific record — a mistake found three steps later costs more than pausing to ask.

1. Ask rather than improvise

Ambiguous parameter? Force field? Stop and ask.

2. Validate before proceeding

Check energy, T , P , RMSD after each stage.

3. Never assume a prior step succeeded

Check exit codes, log files, output files explicitly.

4. Prefer known-good protocols

No new flags without explanation and confirmation.

5. Surface uncertainty explicitly

Protonation state? Box size? Ion conc.? Say so.

6. Errors and warnings are blockers

Understand and resolve — never suppress.



Compute Infrastructure

Host	Role	SLURM	GPU	Claude
local (Mac)	Planning, analysis, writing	—	—	✓
rubisco	Remote workstation	—	yes	✓
villin	Remote workstation	—	yes	✓
insulin	GPU server	yes	yes	✓
calmodulin	GPU server (HP)	yes	yes	✓
agamede	EHU HPC cluster	yes	yes	×
hyperion	DIPC HPC cluster	yes	yes	×

/gh-propagate syncs ~/ .claude to all ✓ machines via git

HPC clusters receive SLURM scripts prepared on interactive machines



Standard Project Layout

```
<ProjectName>/
|-- data/                # MD output (NOT
    git)
|   |-- *.xtc   *.gro   *.edr
|   '-- *.cpt   *.tpr   *.log
|-- analysis/
|   |-- notebooks/      # Jupyter (git)
|   '-- scripts/        # Python  (git)
|-- setup/              # mdp, top, itp,
    gro
|-- scripts/            # SLURM scripts
|-- README.md
'-- .gitignore
```

Key features

- data/ excluded from git — large binary files
- Analysis code tracked for reproducibility
- Same layout on *all* machines — predictable paths
- new_project.sh scaffolds the tree and makes an initial commit

```
# Create a new project:
~/Research/Projects/new_project.sh
\
  MyProjectName
```




Skills Library: BioKT Collection

Category	Skill	What it does
Simulation prep	/gmx-prep	GROMACS: EM → NVT → NPT → production
	/gmx-install	Build GROMACS + PLUMED from source
Enhanced sampling	/plumed-us	Umbrella sampling + WHAM
	/plumed-metad	Well-tempered metadynamics + sum_hills
	/plumed-remd	US-REMD, bias-exchange, PT-METAD
Parameterization	/amber-parameterize	SMILES → GAFF2 → GROMACS (.itp)
Quantum chem.	/dft-homo-lumo	PySCF DFT → HOMO/LUMO energies
Infrastructure	/jobs	Track simulation jobs in JOBS.md
	/sim-paths	SSH pre-flight file check
	/bash-parallelize	Parallel subagents for independent loops
Knowledge mgmt	/overview	Aggregate TODO/CONTEXT across machines
	/ctxt	Save/restore session context
	/bitacora	Append to lab notebook
	/tyler	PDF papers → markdown wiki



Workflow stages



Default settings

Force field: ff99SBws-STQ

Water model: TIP4P/2005

Salt: 150 mM NaCl

$T = 300$ K, $P = 1$ bar

Timestep: 2 fs; cutoffs: 1.0 nm

Electrostatics: PME

Advanced options

AWH (free energy along a CV)

H-REMD (Hamiltonian replica exchange)

SLURM submission templates included

Ligand integration via ACPYPE



Enhanced Sampling with PLUMED

/plumed-us

Umbrella Sampling

Harmonic restraint windows
Serial or MPI/REMD execution
WHAM free energy profile
Bootstrap error estimation
Convergence check (7 stages)

/plumed-metad

Well-Tempered MetaD

Adaptive Gaussian hills
Biasfactor guidance
sum_hills FES
Trajectory reweighting
Block error analysis


/plumed-remd

Replica Exchange

US-REMD, bias-exchange
PT-METAD temp. ladders
gmx_mpi -multidir
Trajectory demultiplexing
WHAM with weights

All three skills produce validated `plumed.dat` files, mdp modifications, and SLURM scripts.





Key choices

Force field: **GAFF2** (General AMBER FF)

Charges: **AM1-BCC** (default) or RESP

Python wrapper automates the full pipeline

Output & integration

Output: `LIG.itp`, `LIG.gro` (GROMACS-ready)


Integrates directly with `/gmx-prep`

DFT-optimized geometries supported as input

RDKit or Open Babel for initial 3D generation



Pipeline



Settings

Functional: CAM-B3LYP

alts: B3LYP, wB97X-D3, PBEo

Basis set: def2-TZVP

high-acc: def2-QZVP

Convergence: 10^{-6} Hartree

Grid: level 4

Integration

DFT-optimized geometry →

/amber-parameterize

RESP charges from DFT wavefunction

Useful for metal-binding residues and cofactors



Workflow Infrastructure

/jobs

Maintains JOBS.md — a table of all running and completed simulations across every machine.

- Job name, machine, SLURM ID
- Start date, status, notes
- Archive completed jobs

/sim-paths

SSH pre-flight check before job submission.

- Verifies all required files exist on the target machine
- Checks .gro, .top, .mdp, .dat, .sh
- Catches missing files before queue wait

/bash-parallelize

Spawns parallel subagents for independent tasks.

- Detects independence in bash loops
- Load-aware concurrency control
- Collects and reports results

Example: analyse 8 replicas in parallel instead of sequentially with a single /bash-parallelize call.



Knowledge & Session Management

`/overview`

Aggregates `CONTEXT.md` and `TODO.md` from all interactive machines via SSH. Produces a prioritised action plan for the day or week. Run from the local Mac.

`/ctxt save / load`

Checkpoints the current session state to `CONTEXT.md`. Run before `/clear` to avoid losing context. Restored automatically at next session start if the file is present in the project directory.

`/bitacora`


Appends a structured entry to `LABNOTEBOOK.md`: summary, decisions made, files touched, figures, next steps. Run before `/clear` to capture the session's scientific record.

`/tyler`

Converts a folder of PDF papers into a two-tier markdown wiki: an `index.md` (~400 tokens/paper) and per-paper papers/`*.md` files read on demand. Keeps literature accessible within Claude's context window.



A Day in the Lab



Morning: /overview reveals what is running and what needs attention

Setup: Claude writes mdp files, prepares topology, generates SLURM script

Submit: /diy executes proposed commands; /jobs records the new entry

Analyse: notebooks with MDtraj, numpy, pandas — Claude helps debug and extend

End of session: /bitacora captures decisions; /ctxt save; then /clear



What You Can Do Starting Monday

First steps

- Install: `curl -fsSL https://claude.ai/install.sh | bash`
- Write your `CLAUDE.md` — role, standards, infrastructure
- Run `/init` in one project to get a local `CLAUDE.md`
- Try `/dft-homo-lumo` on a molecule you know well

Build from there

- Use `/gmx-prep` for your next simulation setup
- Run `/overview` each morning to track active work
- End sessions with `/bitacora` and `/ctxt save`
- Add skills as your workflows mature

Start with tasks you already know how to do. That way you can tell when Claude gets it right — and when it does not.



A Caution Before You Start

The real risk

- Not that Claude will fail spectacularly
- But that it will *succeed just enough* to prevent you from developing your own understanding
- A researcher who cannot evaluate the output is not doing research — they are transcribing an AI
- Deep intuitions about your system take years to build; they cannot be borrowed

The rule

If you cannot tell whether Claude's output is right, you are not yet in a position to trust it.

Use Claude to go *faster* once you know where you are going.

Do not use it to avoid learning the map.



Further Reading & Resources

Documentation

- **Anthropic Claude Code docs**

`https:`
`//code.claude.com/docs`

- **Claude Blattman**

`https:`
`//claudeblattman.com`

Broader perspective

- **One Useful Thing** —
Ethan Mollick

`https:`
`//www.oneusefulthing.org`

- **Tyler Cowen** — AI and
expertise

`https://www.youtube.com/`
`watch?v=aJlg6o0A_Js`

YouTube Tutorials

- John Kim
- Nick Saraev
- ... *and many others*

